

YASH KATIYAR

+91-9129686436 | yashkatiyar2405@gmail.com | github.com/Yashkatiyar24 | LinkedIn | Portfolio |

SUMMARY

AI/ML Engineer with production experience building LangChain agents, RAG pipelines, and LLM-powered SaaS applications. Deployed 3 AI systems handling 1,000+ inference requests. Proficient in Python 3.10+, FastAPI, LangGraph, vector databases (FAISS, Pinecone, ChromaDB), and LLMOps across OpenAI, Gemini, and Claude APIs. Experienced in async workflows, streaming inference pipelines, and scalable deployment on AWS and Vercel.

EXPERIENCE

App Developer Intern — Web Winners Tech Solution, Chandigarh (Remote)

Feb 2026 – Present

- Built enterprise ERP PHP with 500+ RBAC and workflow automation modules, processing 500+ daily transactions.
- Optimized 25+ DB queries via indexing and restructuring, reducing avg response time by **40%** under 200+ concurrent users.
- Architected RESTful API layer connecting 10+ modules, automating 15+ manual workflows and cutting integration time by 30%.

Full-Stack AI Developer — EZ Care AI, France (Remote)

Dec 2025 – Feb 2026

- Built LLM-driven wellness platform (Next.js, TypeScript, Supabase) serving **200+ beta users** via async inference pipelines.
- Engineered EZBuddy AI assistant generating **1,000+ personalized** meal/fitness plans using prompt-engineered LLM inference on real-time user data.

PROJECTS

Real-Time RAG Web AI Assistant

Nov – Dec 2025

Python, LangChain, LlamaIndex, FastAPI, FAISS, OpenAI, Ollama, Streamlit [Live]

- Built LangChain RAG agent with DuckDuckGo retrieval + FAISS vector store; achieved **sub-5s** latency and **92% relevance accuracy** over 100+ test queries with dual LLM backends (OpenAI, Ollama).
- Deployed async FastAPI + Streamlit app with real-time streaming responses, automated citation tracking, and observability hooks for factual traceability.

AI Coding Agent

April 2026

Python, LangGraph, LangChain, asyncio [GitHub]

- Built LangGraph-orchestrated coding agent with stateful multi-step reasoning (plan → code → test → fix), leveraging async tool-use loops and structured state-machine execution.
- Exposed agent via **async FastAPI streaming endpoints**; implemented prompt-level cost optimization reducing token usage ~30% via selective context injection.

ColdClaw — AI Cold Email Generator SaaS

Mar 2026

Next.js, Firebase, Gemini API, Prompt Engineering [Live]

- Built AI SaaS generating full email sequences in **under 60s** with Gemini API and 5 tone presets; produced **100+ campaigns** with zero duplication.
- Engineered dynamic prompt system with feedback-based personalization achieving **95%+ output accuracy**; applied LLMops cost controls via prompt caching and token budgeting.

ACHIEVEMENTS

- 1st Place** — 24-Hour iOS Hackathon (2026): Shipped SwiftUI app (MVVM) in 24 hours, beating 50+ teams.
- Top 5** — Odoo Hackathon, IIT Madras: Designed enterprise ERP workflow outperforming 100+ teams.
- International** — Hankao University, Beijing: Enhanced result platform actively used by 5,000+ students.

SKILLS

Agent & LLM: LangChain, LangGraph, LlamaIndex, OpenAI API, Gemini API, Claude API, Ollama, RLHF

LLMOps & Vector DBs: Prompt Engineering, Cost Optimization, Token Budgeting, Streaming, Observability, FAISS, Pinecone, ChromaDB, Redis

AI/ML & Languages: PyTorch, TensorFlow, Hugging Face Transformers, Scikit-learn, BERT, ViLT, Python, TypeScript, SQL, Git

EDUCATION

KIET Group of Institutions, Ghaziabad — B.Tech, CSE (Artificial Intelligence)

2023–2027

IIT Madras — B.S. in Data Science

2023–2027